

# Principles of Robot Autonomy I

Motion planning I: graph search algorithms

# Attendance Form





## OPEN HOUSE

October 16th, 2025 @ 4pm

Durand Building, Room 023

Food provided!

Zoom link:

<https://stanford.zoom.us/j/97988116208?pwd=Geb3aFdoFsTL9J97GOX4XoNpsXPoMn.1>

### Schedule


<b>4PM</b>	Introduction by Professor Marco Pavone
<b>4:05PM</b>	<p>5 minute lightning talks about the lab's research directions and applications</p> <ul style="list-style-type: none"><li>• <i>Foundation Models for Next-Generation Autonomy Stacks</i></li><li>• <i>Test-Time Scaling and Reasoning for Robotics</i></li><li>• <i>Physical AI Safety: Monitoring, Alignment, and Guardrails</i></li><li>• <i>Data Flywheels and Data Attribution</i></li><li>• <i>Blending AI and Optimization/Control</i></li><li>• <i>Application domains: Space Robotics, Manipulators, Quadrupeds, and more</i></li></ul>
<b>4:35PM</b>	<p>Open discussion. Opportunity to ask questions about the lab, specific research directions, classes, research experience, or anything else.</p> <p>If time permits, we can include a tour of the lab and the Space Robotics Facility.</p>

# Agenda

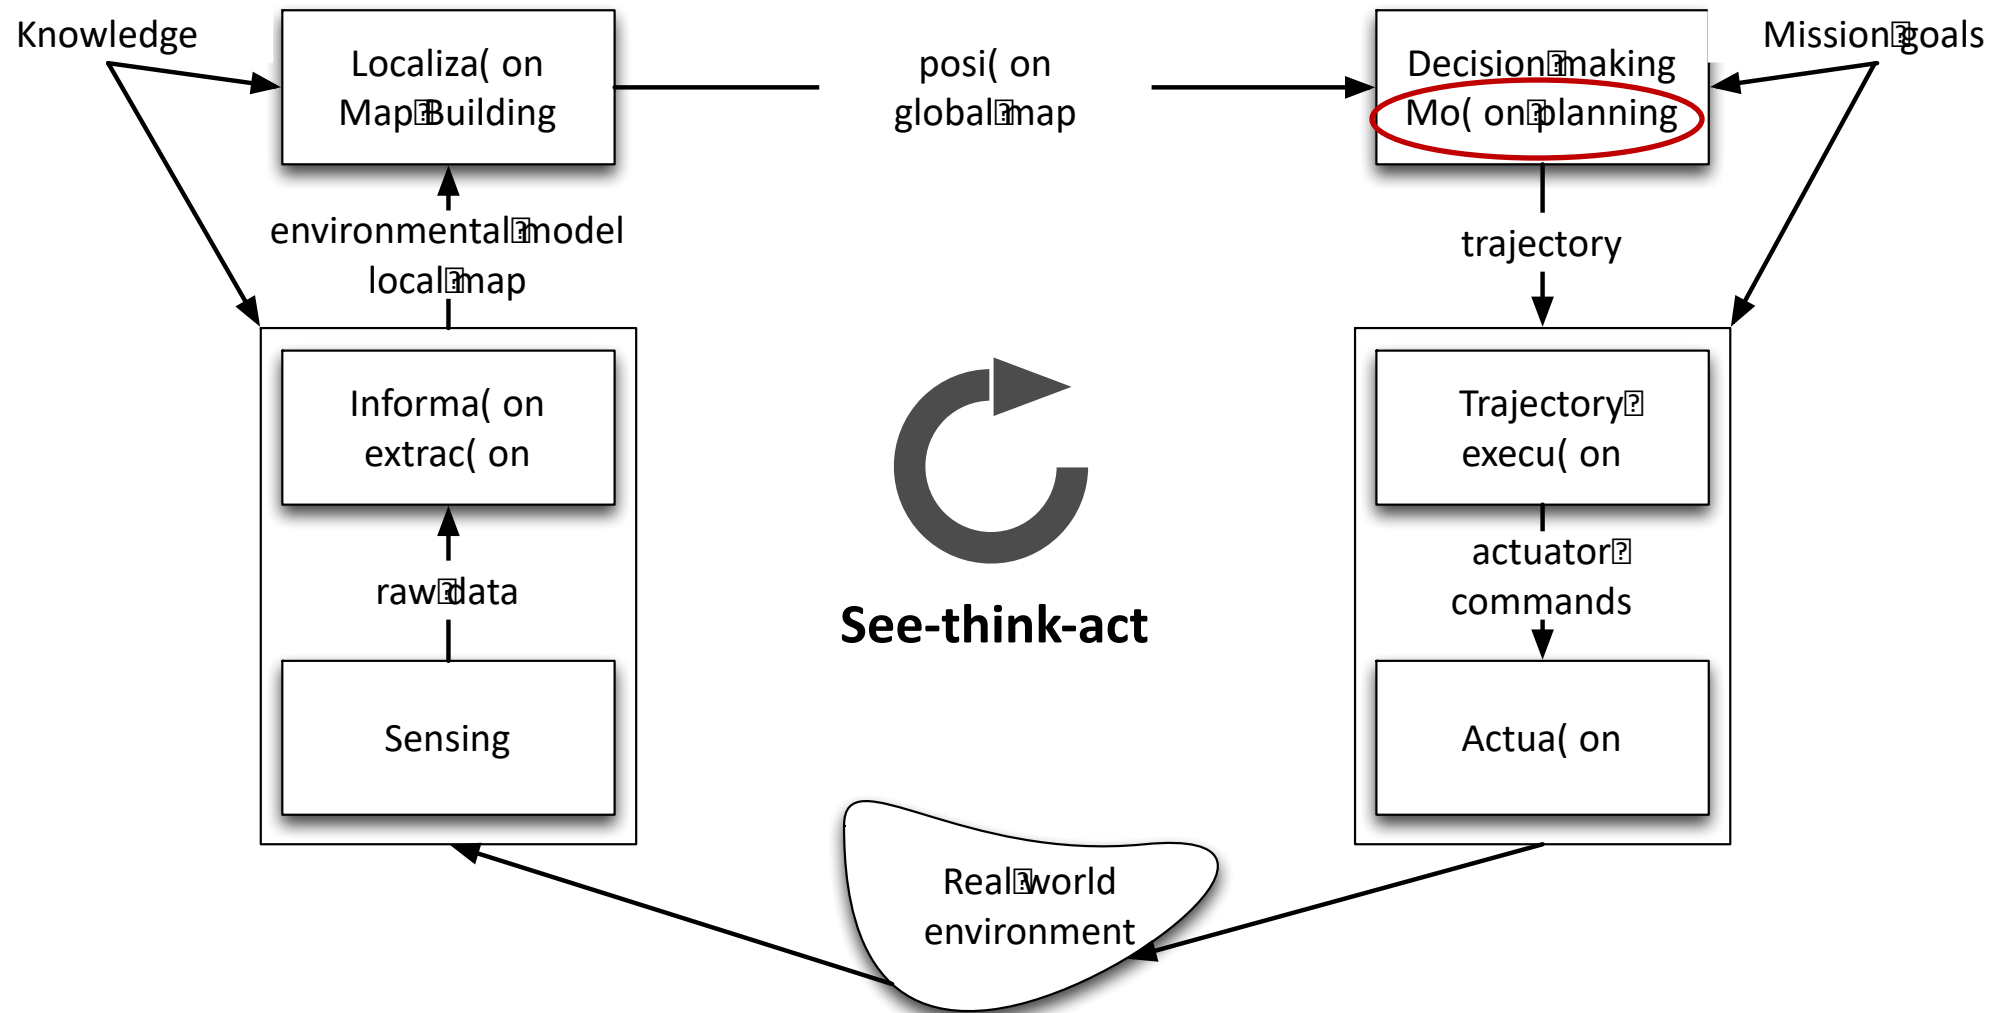
- Agenda
  - Introduction to motion planning
  - Search-based algorithms for motion planning
  - Configuration spaces and combinatorial motion planning
- Readings:
  - Chapter 4, sections 4.1 – 4.2 in D. Gammelli, J. Lorenzetti, K. Luo, G. Zardini, M. Pavone. *Principles of Robot Autonomy*. 2026.


# Motion planning

**Problem definition:** Compute sequence of actions that drives a robot from an initial condition to a terminal condition while avoiding obstacles, respecting motion constraints, and *possibly* optimizing a cost function



# The see-think-act cycle








8x

# More examples of motion planning

- Steering autonomous vehicles
- Controlling humanoid robot
- Surgery planning
- Protein folding
- ...







# Some history

- Formally defined in the 1970s
- Development of exact, combinatorial solutions in the 1980s
- Development of sampling-based methods in the 1990s
- Deployment on real-time systems in the 2000s
- Current research: inclusion of differential and logical constraints, planning under uncertainty, parallel implementation, and more

# Simplest setup

- Assume 2D workspace:  $\mathcal{W} \subseteq \mathbb{R}^2$
- $\mathcal{O} \subset \mathcal{W}$  is the obstacle region with polygonal boundary
- Robot is a rigid polygon
- **Problem:** given initial placement of robot, compute how to gradually move it into a desired goal placement so that it never touches the obstacle region




# Popular approaches

- *Potential fields* [Rimon, Koditschek, '92]: create forces on the robot that pull it toward the goal and push it away from obstacles
- *Grid-based planning* [Stentz, '94]: discretizes problem into grid and runs a graph-search algorithm (Dijkstra, A\*, ...)
- *Combinatorial planning* [LaValle, '06]: constructs structures in the configuration (C-) space that completely capture all information needed for planning
- *Sampling-based planning* [Kavraki et al, '96; LaValle, Kuffner, '06, etc.]: uses collision detection algorithms to probe and incrementally search the C-space for a solution, rather than completely characterizing all of the  $C_{\text{free}}$  structure

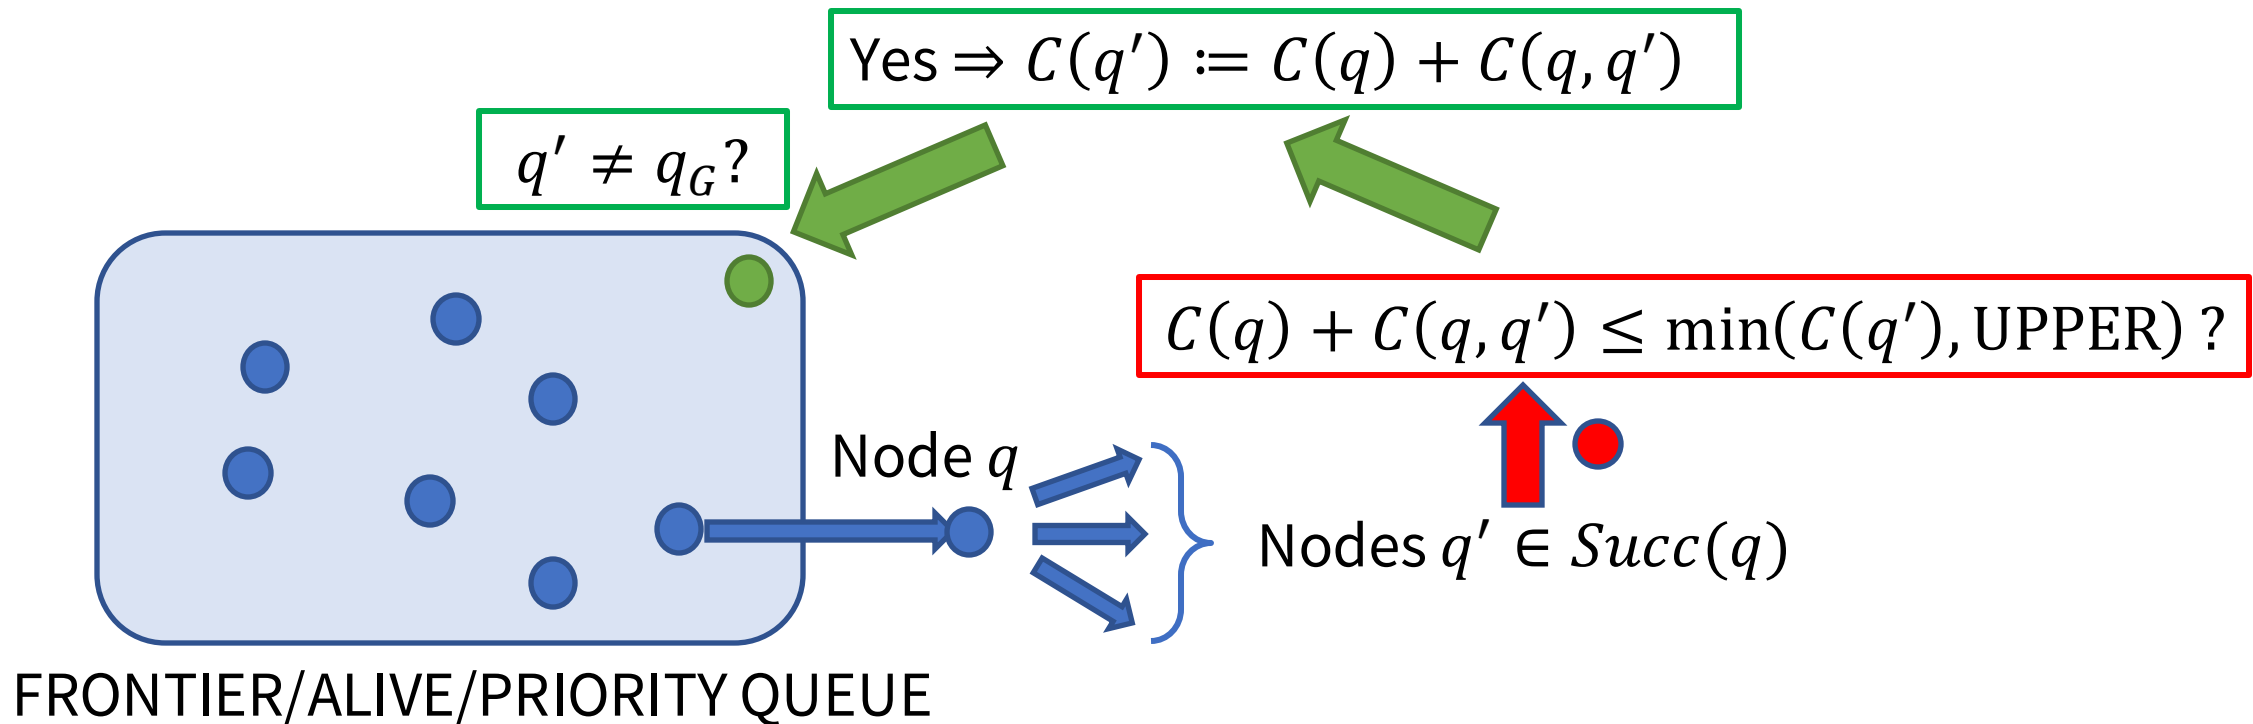
# Grid-based approaches

- Discretize the continuous world into a grid
  - Each grid cell is either free or forbidden
  - Robot moves between adjacent free cells
  - **Goal:** find sequence of free cells from start to goal
- Mathematically, this corresponds to pathfinding in a discrete graph  $G = (V, E)$ 
  - Each vertex  $v \in V$  represents a free cell
  - Edges  $(v, u) \in E$  connect adjacent grid cells



# Graph search algorithms

- Having determined decomposition, how to find “best” path?
- **Label-Correcting Algorithms:**  $C(q)$ : *cost-of-arrival* from  $q_I$  to  $q$



# Label correcting algorithm

**Step 1.** Remove a node  $q$  from frontier queue and for each child  $q'$  of  $q$ , execute step 2

**Step 2.** If  $C(q) + C(q, q') \leq \min(C(q'), \text{UPPER})$ , set  $C(q') := C(q) + C(q, q')$  and set  $q$  to be the parent of  $q'$ . In addition, if  $q' \neq q_G$ , place  $q'$  in the frontier queue if it is not already there, while if  $q' = q_G$ , set UPPER to the new value  $C(q) + C(q, q_G)$


**Step 3.** If the frontier queue is empty, terminate, else go to step 1

**Initialization:** set the labels of all nodes to  $\infty$ , except for the label of the origin node, which is set to 0

# GetNext() ?


**Depth-First-Search (DFS):** Maintain  $Q$  as a **stack** – Last in/first out

- Lower memory requirement (only need to store part of graph)



**Breadth-First-Search (BFS, Bellman-Ford):** Maintain  $Q$  as a **list** – First in/first first out

- Update cost for all edges up to current depth before proceeding to greater depth
- Can deal with negative edge (transition) costs





**Best-First (BF, Dijkstra):** Greedily select next  $q$ :  $q = \operatorname{argmin}_{q \in Q} \mathcal{C}(q)$

- Node will enter the frontier queue at most *once*
- Requires costs to be non-negative

# Correctness and improvements

## Theorem

If a feasible path exists from  $q_I$  to  $q_G$ , then algorithm terminates in finite time with  $C(q_G)$  equal to the optimal cost of traversal,  $C^*(q_G)$ .



\* [https://en.wikipedia.org/wiki/Dijkstra%27s\\_algorithm](https://en.wikipedia.org/wiki/Dijkstra%27s_algorithm)



# A\*: Improving Dijkstra

- Dijkstra orders by optimal “*cost-to-arrival*”
- Faster results if order by “*cost-to-arrival*”+ (approximate) “*cost-to-go*”
- That is, strengthen test

$$C(q) + C(q, q') \leq \text{UPPER}$$

to

$$C(q) + C(q, q') + h(q') \leq \text{UPPER}$$

where  $h(q)$  is a heuristic for optimal cost-to-go (specifically, a positive *underestimate*)


- In this way, fewer nodes will be placed in the frontier queue
- This modification still guarantees that the algorithm will terminate with a shortest path
- Many variations are possible... see (Problem 2 in pset 2)

# Grid-based approaches: summary

- Pros:
  - Simple and easy to use
  - Fast (for some problems)
- Cons:
  - Resolution dependent
    - Not guaranteed to find solution if grid resolution is not small enough
  - Limited to simple robots
    - Grid size is exponential in the number of DOFs


# Back to continuous motion planning

- Assume 2D workspace:  $\mathcal{W} \subseteq \mathbb{R}^2$
- $\mathcal{O} \subset \mathcal{W}$  is the obstacle region with polygonal boundary
- Robot is a rigid polygon
- **Problem:** Given initial placement of robot, compute how to gradually move it into a desired goal placement so that it never touches the obstacle region






# Back to continuous motion planning

**Key point:** motion planning problem described in the real-world, but it really lives in another space -- the **configuration** (C-) space!




# Configuration space


- **C- space:** captures all degrees of freedom (all rigid body transformations)
- More in detail, let  $\mathcal{R} \subset \mathbb{R}^2$  be a polygonal robot (e.g., a triangle)
- The robot can rotate by angle  $\theta$  or translate  $(x_t, y_t) \in \mathbb{R}^2$
- Every combination  $q = (x_t, y_t, \theta)$  yields a *unique* robot placement: **configuration**
- So C- space is a subset of  $\mathbb{R}^3$
- Note:  $\theta \pm 2\pi$  yields equivalent rotations  $\Rightarrow$  C- space is:  $\mathbb{R}^2 \times \mathcal{S}^1$
- Concept of C- space extends naturally to higher dimensions (e.g., robot linkages)



# Configuration free space

- The subset  $\mathcal{F} \subseteq \mathcal{C}$  of all collision free configurations is the **free space**






**Bottom line:** *explicitly* computing  $C$  free spaces in high-dimensional settings is hard!




# Planning in $C$ -space

- Let  $R(q) \subset W$  denote set of points in the world occupied by robot when in configuration  $q$
- Robot in collision  $\Leftrightarrow R(q) \cap O \neq \emptyset$
- Accordingly, *free space* is defined as:  $C_{\text{free}} = \{q \in C \mid R(q) \cap O = \emptyset\}$
- Path planning problem in  $C$ -space: compute a **continuous** path:  $\tau: [0,1] \rightarrow C_{\text{free}}$ , with  $\tau(0) = q_I$  and  $\tau(1) = q_G$



# Combinatorial planning

- Combinatorial approaches to motion planning find paths through continuous configuration space **without resorting to approximations**
- **Key idea:** compute a roadmap, which provides a discrete representation of continuous motion planning problem without losing any of the original connectivity information needed to solve it
- Such approaches are typically complete (i.e., guaranteed to find a solution), but are typically limited to **small number of DOFs** due to the challenge of *exactly* computing  $C$  free spaces



A roadmap is a graph in which each vertex is a configuration in  $C_{\text{free}}$  and each edge is a path through  $C_{\text{free}}$  that connects a pair of vertices

# Next time: sampling-based planning

